

Chapter 2 Lab Practices: Integer Representations, Arithmetic, Floating-Point, and Additional Bit Manipulation

DS2040, Satyajit Das

January 16, 2025

Introduction

This lab collection covers hands-on programming tasks related to:

- Integer representations (two's complement, overflow, sign extension)
- Integer arithmetic (truncation, division, multiplication)
- Floating-point representation and arithmetic (precision, rounding, Inf, NaN)
- Bit manipulation (shifts, masking, rotating bits, setting/clearing bits)

By completing these programs, students will gain a deeper understanding of how data is represented and manipulated at the machine level.

1 Programs

1. Printing Integer Representations

Purpose: Demonstrate how integers can be displayed in different bases (decimal, octal, hexadecimal).

Listing 1: print_int_representations.c

```
1 #include <stdio.h>
2
3 int main(void) {
4     int num = 100;
5
6     printf("Decimal: %d\n", num);
7     printf("Octal: %o\n", num);
8     printf("Hexadecimal: %x\n", num);
9
10    return 0;
11 }
```

Key Concepts:

- %d for decimal
- %o for octal
- %x for hexadecimal

2. Two's Complement Observation

Purpose: Observe how two's complement wraps around when dealing with signed integers.

Listing 2: twos_complement_wrap.c

```

1 #include <stdio.h>
2 #include <limits.h>
3
4 int main(void) {
5     int x = INT_MAX;
6     printf("INT_MAX = %d\n", x);
7     x = x + 1; // Overflow on a 2's complement system
8     printf("After adding 1 to INT_MAX: %d\n", x);
9
10    int y = INT_MIN;
11    printf("INT_MIN = %d\n", y);
12    y = y - 1; // Underflow
13    printf("After subtracting 1 from INT_MIN: %d\n", y);
14
15    return 0;
16 }
```

Key Concepts:

- Two's complement overflow/underflow behavior
- INT_MAX and INT_MIN

3. Sign Extension in Arithmetic

Purpose: Demonstrate what happens when a smaller signed integer is promoted to a larger type.

Listing 3: sign_extension.c

```

1 #include <stdio.h>
2 #include <stdint.h>
3
4 int main(void) {
5     int16_t smallNum = -12345; // 16-bit signed
6     int32_t largeNum = smallNum; // Implicit sign extension
7
8     printf("16-bit smallNum = %d\n", smallNum);
9     printf("32-bit largeNum (extended) = %d\n", largeNum);
10
11    return 0;
12 }
```

Key Concepts:

- Sign extension from smaller to larger type
- Preserving negative values

4. Bitwise Operations on Integers

Purpose: Review AND, OR, XOR, bit-shifts, and see how they affect integer values.

Listing 4: integer_bitwise_ops.c

```
1 #include <stdio.h>
2 #include <stdint.h>
3
4 int main(void) {
5     uint32_t a = 0xF0F0F0F0;
6     uint32_t b = 0xAAAA5555;
7
8     printf("a = 0x%08X\n", a);
9     printf("b = 0x%08X\n", b);
10
11     printf("a & b = 0x%08X\n", (a & b));
12     printf("a | b = 0x%08X\n", (a | b));
13     printf("a ^ b = 0x%08X\n", (a ^ b));
14     printf("a << 4 = 0x%08X\n", (a << 4));
15     printf("b >> 4 = 0x%08X\n", (b >> 4));
16
17     return 0;
18 }
```

Key Concepts:

- Bitwise &, |, ^
- Left and right shifting (<<, >>)
- Hexadecimal display format (%08X)

5. Endianness Check

Purpose: Illustrates how to check if a system is little-endian or big-endian using a small integer.

Listing 5: endianness_check.c

```
1 #include <stdio.h>
2
3 int main(void) {
4     unsigned int x = 0x12345678;
5     unsigned char *ptr = (unsigned char *)&x;
6
7     printf("Memory representation: ");
8     for(int i = 0; i < 4; i++) {
9         printf("%02X ", ptr[i]);
10     }
```

```

10     }
11     printf("\n");
12
13     if (ptr[0] == 0x78) {
14         printf("System is Little-Endian.\n");
15     } else {
16         printf("System is Big-Endian.\n");
17     }
18
19     return 0;
20 }

```

Key Concepts:

- Interpreting memory byte by byte
- Detecting endianness by examining the least significant byte

6. Unsigned Arithmetic Overflow

Purpose: Show what happens when you exceed the maximum of an unsigned type.

Listing 6: unsigned_overflow.c

```

1  #include <stdio.h>
2  #include <limits.h>
3
4  int main(void) {
5      unsigned int max = UINT_MAX;
6      printf("UINT_MAX = %u\n", max);
7
8      unsigned int overflowed = max + 1;
9      printf("After adding 1 to UINT_MAX: %u\n", overflowed);
10
11     return 0;
12 }

```

Key Concepts:

- Unsigned integer range
- Wrap-around behavior (mod 2^n)

7. Multiplication and Division Truncation

Purpose: Demonstrate how integer multiplication and division behave with truncation and overflow.

Listing 7: int_mult_div.c

```

1  #include <stdio.h>
2  #include <limits.h>
3
4  int main(void) {
5      int a = 20000;

```

```

6     int b = 20000;
7     long long result = (long long)a * b; // to see exact product
8
9     printf("a * b (int) = %d (overflow expected)\n", a * b);
10    printf("a * b (long long) = %lld\n", result);
11
12    int numerator = 7;
13    int denominator = 2;
14    printf("Integer division of 7 / 2 = %d\n", numerator / denominator);
15
16    return 0;
17 }

```

Key Concepts:

- Overflow when int is too small
- Truncation in integer division

8. Fixed-Width Integer Arithmetic

Purpose: Explore arithmetic with types like `int8_t`, `int16_t`, `int32_t`, which can overflow more easily.

Listing 8: `fixed_width_integers.c`

```

1  #include <stdio.h>
2  #include <stdint.h>
3
4  int main(void) {
5      int8_t x8 = 120; // near INT8_MAX is 127
6      int16_t x16 = 30000; // near INT16_MAX is 32767
7
8      printf("x8 = %d\n", x8);
9      x8 = x8 + 10; // Overflow might occur
10     printf("After adding 10: x8 = %d\n", x8);
11
12     printf("x16 = %d\n", x16);
13     x16 = x16 + 5000; // Might overflow
14     printf("After adding 5000: x16 = %d\n", x16);
15
16     return 0;
17 }

```

Key Concepts:

- Fixed-width integers from `<stdint.h>`
- Overflow in smaller data types

9. Simple Floating-Point Representation

Purpose: Show how floats store approximate values and compare differences with doubles.

Listing 9: float_rep.c

```
1 #include <stdio.h>
2
3 int main(void) {
4     float f = 3.141592653589793f; // single precision
5     double d = 3.141592653589793; // double precision
6
7     printf("float f: %.9f\n", f);
8     printf("double d: %.15f\n", d);
9
10    return 0;
11 }
```

Key Concepts:

- Single precision vs. double precision
- Displaying more decimal places than the type can store accurately

10. Rounding Errors in Floating-Point Arithmetic

Purpose: Demonstrate how floating-point addition/subtraction leads to rounding errors.

Listing 10: float_rounding.c

```
1 #include <stdio.h>
2
3 int main(void) {
4     float sum = 0.0f;
5     for(int i = 0; i < 1000; i++) {
6         sum += 0.01f;
7     }
8     printf("Expected 10.0, got %f\n", sum);
9
10    return 0;
11 }
```

Key Concepts:

- Floating-point rounding errors
- Accumulated precision loss

11. Casting Between Floating-Point and Integer

Purpose: Illustrate truncation and rounding when casting between float/double and integers.

Listing 11: float_int_cast.c

```
1 #include <stdio.h>
2
3 int main(void) {
4     float f = 3.99f;
5     int i = (int)f; // cast to int truncates toward zero
```

```

6
7     printf("float f = %.2f\n", f);
8     printf("After casting to int: i = %d\n", i);
9
10    double d = 123456789.0;
11    int j = (int)d; // might exceed int range if too large
12    printf("double d = %.0f\n", d);
13    printf("After casting to int: j = %d\n", j);
14
15    return 0;
16 }

```

Key Concepts:

- Truncation from float/double to int
- Potential overflow if float/double is too large

12. Exploring printf Format Specifiers for Floating-Point

Purpose: Practice printing floating-point numbers in various formats.

Listing 12: float_format_specifiers.c

```

1 #include <stdio.h>
2
3 int main(void) {
4     double val = 12345.6789;
5
6     printf("Default: %f\n", val);
7     printf("Fixed decimal (2 digits): %.2f\n", val);
8     printf("Scientific notation: %e\n", val);
9     printf("Hexadecimal notation: %a\n", val);
10
11    return 0;
12 }

```

Key Concepts:

- %f, %e, %a format specifiers
- Scientific and hexadecimal floating-point representation

13. Floating-Point Comparisons Pitfall

Purpose: Demonstrate why direct equality checks on floating-point numbers are problematic.

Listing 13: float_comparison.c

```

1 #include <stdio.h>
2 #include <math.h>
3
4 int main(void) {
5     double x = 0.1 + 0.2;

```

```

6     double y = 0.3;
7
8     if (x == y) {
9         printf("x == y\n");
10    } else {
11        printf("x != y (due to floating-point precision)\n");
12    }
13
14    // A better way:
15    if (fabs(x - y) < 1e-9) {
16        printf("x and y are close enough\n");
17    } else {
18        printf("x and y differ by more than 1e-9\n");
19    }
20
21    return 0;
22 }

```

Key Concepts:

- Precision errors in floating-point arithmetic
- Using an epsilon for comparisons

14. Checking Floating-Point Range and Infinity

Purpose: Understand how floating-point numbers handle very large values or division by zero.

Listing 14: float_infinity.c

```

1  #include <stdio.h>
2  #include <math.h>
3
4  int main(void) {
5      double huge = 1.0e308;
6      double bigger = huge * 1000.0; // Likely overflow to infinity
7
8      printf("huge = %e\n", huge);
9      printf("bigger = %e\n", bigger);
10
11     double zero = 0.0;
12     double inf = 1.0 / zero;
13     printf("1.0 / 0.0 = %f\n", inf); // Infinity
14
15     double nanVal = 0.0 / zero;
16     printf("0.0 / 0.0 = %f (NaN)\n", nanVal);
17
18     return 0;
19 }

```

Key Concepts:

- IEEE 754 behavior for overflow (Inf) and invalid operations (NaN)
- Distinguishing finite numbers vs. `inf`, `NaN`

15. Mixing Integer and Floating-Point Arithmetic

Purpose: Examine how C performs type promotion when expressions have both integer and floating types.

Listing 15: mixed_arithmetic.c

```
1 #include <stdio.h>
2
3 int main(void) {
4     int i = 5;
5     float f = 2.5f;
6
7     float result1 = i + f;          // i promoted to float
8     double result2 = (double)i / 2; // i stays int, but result is double
9
10    printf("i + f = %.2f\n", result1);
11    printf("(double)i / 2 = %.2f\n", result2);
12
13    // Another subtlety:
14    int j = (int)(i * f); // i*f is float, then truncated
15    printf("(int)(i * f) = %d\n", j);
16
17    return 0;
18 }
```

Key Concepts:

- Integer-to-float promotion rules
- Truncation vs. rounding

16. Reversing Bits in an Unsigned Integer

Purpose: Practice manipulating bits to reverse the bit order of an unsigned integer.

Listing 16: bit_reverse.c

```
1  #include <stdio.h>
2  #include <stdint.h>
3
4  uint32_t reverse_bits(uint32_t x) {
5      uint32_t reversed = 0;
6      for(int i = 0; i < 32; i++) {
7          reversed <<= 1;
8          reversed |= (x & 1);
9          x >>= 1;
10     }
11     return reversed;
12 }
13
14 int main(void) {
15     uint32_t num = 0x12345678;
16     uint32_t rev = reverse_bits(num);
17
18     printf("Original: 0x%08X\n", num);
19     printf("Reversed: 0x%08X\n", rev);
20
21     return 0;
22 }
```

Key Concepts:

- Looping to shift and accumulate bits
- Bitwise masking (& 1)

17. Rotating Bits Left and Right

Purpose: Demonstrate how to rotate bits around in a 32-bit integer (different from shifting).

Listing 17: bit_rotate.c

```
1  #include <stdio.h>
2  #include <stdint.h>
3
4  uint32_t rotate_left(uint32_t x, int n) {
5      return (x << n) | (x >> (32 - n));
6  }
7
8  uint32_t rotate_right(uint32_t x, int n) {
9      return (x >> n) | (x << (32 - n));
10 }
11
12 int main(void) {
13     uint32_t val = 0xABCDEF12;
```

```

14
15     printf("Original: 0x%08X\n", val);
16     printf("Rotated L4: 0x%08X\n", rotate_left(val, 4));
17     printf("Rotated R4: 0x%08X\n", rotate_right(val, 4));
18
19     return 0;
20 }

```

Key Concepts:

- Difference between rotate and shift (shifts introduce zeros)
- | to combine shifted portions

18. Counting Set Bits in an Integer

Purpose: Show how to count the number of 1-bits in a given integer (a.k.a. the Hamming weight).

Listing 18: count_set_bits.c

```

1  #include <stdio.h>
2  #include <stdint.h>
3
4  int count_set_bits(uint32_t x) {
5      int count = 0;
6      while (x) {
7          count += (x & 1);
8          x >>= 1;
9      }
10     return count;
11 }
12
13 int main(void) {
14     uint32_t num = 0xF0F0F0F0;
15     int bits = count_set_bits(num);
16
17     printf("Number = 0x%08X\n", num);
18     printf("Set bits = %d\n", bits);
19
20     return 0;
21 }

```

Key Concepts:

- Iterative approach to count bits
- Bitwise & 1 and right shift

19. Checking the Sign with Bits (Signed vs. Unsigned)

Purpose: Use bitwise operations to check if a signed integer is negative and compare with an unsigned interpretation.

Listing 19: check_sign.c

```

1  #include <stdio.h>
2  #include <stdint.h>
3
4  int main(void) {
5      int32_t s = -42; // signed
6      uint32_t u = (uint32_t)s; // re-interpret bits as unsigned
7
8      printf("Signed s = %d\n", s);
9      printf("Unsigned u = %u\n", u);
10
11     // Checking sign bit manually (assuming 32-bit two's complement)
12     int sign_bit = (s >> 31) & 1;
13     printf("Sign bit of s: %d\n", sign_bit);
14
15     return 0;
16 }

```

Key Concepts:

- Reinterpreting the same bits as signed vs. unsigned
- Extracting the sign bit with a right shift

20. Clearing and Setting Specific Bits

Purpose: Practice using masks to clear (set to 0) and set (to 1) specific bits within an integer.

Listing 20: bit_set_clear.c

```

1  #include <stdio.h>
2  #include <stdint.h>
3
4  int main(void) {
5      uint32_t x = 0xFF00FF00;
6      printf("Original x = 0x%08X\n", x);
7
8      // Clear the lower 8 bits
9      uint32_t mask_clear = 0xFFFFF00; // 11111111 11111111 11111111 00000000
10     x = x & mask_clear;
11     printf("After clearing lower 8 bits = 0x%08X\n", x);
12
13     // Set bits 8 to 15
14     uint32_t mask_set = 0x0000FF00;
15     x = x | mask_set;
16     printf("After setting bits 8-15 = 0x%08X\n", x);
17
18     return 0;
19 }

```

Key Concepts:

- Constructing masks for specific bits

- `&` to clear bits, `|` to set bits

21. Converting Signed to Unsigned and Observing Behavior

Purpose: Show how assigning a negative signed integer to an unsigned variable can change its interpretation.

Listing 21: `signed_to_unsigned.c`

```

1 #include <stdio.h>
2 #include <stdint.h>
3
4 int main(void) {
5     int32_t negativeVal = -12345;
6     uint32_t convertedVal = (uint32_t)negativeVal;
7
8     printf("signed int: %d\n", negativeVal);
9     printf("unsigned int: %u\n", convertedVal);
10    printf("as hex: 0x%08X\n", convertedVal);
11
12    return 0;
13 }
```

Key Concepts:

- Bitwise re-interpretation
- Large unsigned values representing negative signed integers

22. Extracting Bytes from a 32-bit Integer

Purpose: Show how to extract individual bytes (e.g., for network protocols or endianness checks).

Listing 22: `extract_bytes.c`

```

1 #include <stdio.h>
2 #include <stdint.h>
3
4 int main(void) {
5     uint32_t val = 0x12345678;
6     unsigned char *ptr = (unsigned char *)&val;
7
8     printf("val = 0x%08X\n", val);
9     printf("Byte 0: 0x%02X\n", ptr[0]);
10    printf("Byte 1: 0x%02X\n", ptr[1]);
11    printf("Byte 2: 0x%02X\n", ptr[2]);
12    printf("Byte 3: 0x%02X\n", ptr[3]);
13
14    return 0;
15 }
```

Key Concepts:

- Using a `char *` pointer to access underlying bytes
- Endianness and byte ordering

23. Shifting a Negative Number and Observing Sign Extension

Purpose: Demonstrate arithmetic shift vs. logical shift on a negative number (in C, right shift on signed is implementation-defined, but usually arithmetic).

Listing 23: negative_shift.c

```
1 #include <stdio.h>
2 #include <stdint.h>
3
4 int main(void) {
5     int32_t neg = -1;      // 0xFFFFFFFF in two's complement
6     uint32_t uneg = (uint32_t)neg;
7
8     printf("neg (signed) = %d\n", neg);
9     printf("uneg (unsigned) = %u\n", uneg);
10
11     // Right shifting a signed negative
12     int32_t shiftedNeg = neg >> 1; // typically arithmetic shift
13     printf("neg >> 1 = %d (0x%08X)\n", shiftedNeg, (uint32_t)shiftedNeg);
14
15     // Right shifting the unsigned version (logical shift)
16     uint32_t shiftedUNeg = uneg >> 1;
17     printf("uneg >> 1 = %u (0x%08X)\n", shiftedUNeg, shiftedUNeg);
18
19     return 0;
20 }
```

Key Concepts:

- Arithmetic vs. logical shifts
- Sign extension on right shift of signed negative numbers
- Implementation-defined aspects in C

Compilation and Practice Tips

- **Compile Each Program:**

```
1 gcc filename.c -o program
2 ./program
```

- **Experimentation:**

- Change constants and data types to see different behaviors (e.g., `int16_t` vs. `int32_t`).
- Print intermediate results or debug info to trace calculations.

- **Further Practice:**

- Implement functions to convert integers to binary strings.
- Parse and print IEEE 754 binary components (sign, exponent, mantissa).
- Test corner cases with large values, small values, negative values, or extreme floating-point numbers.